

Sungbin Im

(Jake Im)

CISO at QueryPie, working at the intersection of AI agents, security engineering, access control platforms, and agentic AI reliability research.

PROFILE

Role
CISO, QueryPie

Location
Seoul, Korea

Email
jake@querypie.com

GitHub
github.com/jazdr

LinkedIn
linkedin.com/in/binlim

Research Interests

- Time-Series Anomaly Detection and Behavioral Drift Detection
- AI Agents and Multi-Agent Systems
- Agent Reliability and Runtime Trace Evaluation
- Sequential Decision Making and Reinforcement Learning
- AI Safety and Security in Agent Systems

Education

aSSIST University Expected Aug 2026
M.S. in AI and Big Data, in progress

Chung-Ang University 2007
B.A. in Law

Professional Experience

QueryPie, Inc. 2023 – Present
Chief Information Security Officer (CISO)

- Lead information security strategy and governance for an AI-based IT startup.
- Oversee security architecture and risk management across AI platform and access control platform domains.
- Design and manage security-related systems including AI red teaming, guardrails, and access control architecture.
- Contribute to AI-driven platform design across database, system, Kubernetes, and web application environments.
- Design and manage AI security and governance practices, including ISO/IEC 42001-related initiatives.

Selected Technical Writing

QueryPie's New Standard for Penetration Testing

Nov 2024

Lead-authored QueryPie whitepaper on penetration testing methodology and security validation.

QueryPie's DevSecOps Pipeline: Proven to Enhance Development Speed and Stability

Nov 2024

Lead-authored QueryPie whitepaper on DevSecOps automation, vulnerability management, and secure delivery pipelines.

Research Experience

Master's Thesis, in progress

A Trace-Grounded Method for Evaluating Reliability Gates in Agentic Software Workflows

- Decompose agentic workflow reliability into four boundary decisions: acceptance, repair, continuation, and persistence.
 - Design a runtime evidence contract (policy_event_trace) recording gate status, evidence references, and persistence outcomes per run.
 - Evaluate on 240 executed sandbox runs with a 120-row human-reviewed subset and blind second-annotator check (Cohen's kappa = 0.87-0.93).
 - Show that 34 of 105 committed runs were premature commitments: boundary failures invisible to completion-only evaluation.
-

Proposed Doctoral Research Direction

- Detect gradual behavioral drift in agentic AI systems before it becomes observable failure.
 - Model agent decision logs and runtime traces as multivariate time series.
 - Connect anomaly detection with sequential decision-making and policy-level behavioral analysis.
 - Develop practical reliability auditing methods for enterprise AI agents, grounded in security guardrails and governance practice.
-