CV

Sungbin Im (Jake Im)

CISO, QueryPie

Email: jake@querypie.com

Mobile: +82-10-9045-6992

Location: Seoul, Korea

Research Interests

- AI Agents and Multi-Agent Systems
- Agent Orchestration and Decision Structure
- Sequential Decision Making
- Reinforcement Learning for Policy Learning
- AI Safety and Security in Agent Systems

Education

aSSIST Seoul School of Integrated Sciences & Technologies

M.S. in AI and Big Data, in progress

Expected graduation: August 2026

Chung-Ang University

B.A. in Law (Transferred)

Graduated: 2007

Professional Experience

QueryPie, Inc.

Chief Information Security Officer (CISO)

2023 – Present

- Lead information security strategy and governance for an AI-based IT startup
- Oversee security architecture and risk management across AI Platform and Access Control Platform

- Apply AI and multi-agent systems to real-world problem solving and operational risk management
- Design and manage security-related systems including AI Red Teaming, Guardrails, and access control architecture
- Contribute to AI-driven platform design in areas including database, system, Kubernetes, and web application environments
- Design and manage AI security and governance practices, including AI Red Teaming, Guardrails, and ISO/IEC 42001-related initiatives

Research Experience

Master's Thesis (in progress)

Design and Evaluation of a High-Level Orchestrator for Long-Horizon Interactive Tasks

- Study the role of high-level orchestration in hierarchical multi-agent systems
- Analyze orchestration trajectories and decision patterns from real system logs
- Define failure types such as looping, verification failure, and termination errors
- Evaluate performance-cost trade-offs in long-horizon interactive tasks
- Focus on agent selection, ordering, verification timing, and termination policy
- Develop evaluation harnesses for AI agent workflows and experimentation infrastructure for multi-agent performance and failure analysis

Proposed Doctoral Research Direction

- Formalize agent orchestration as a sequential decision problem
- Model selection, ordering, and termination policies of high-level orchestrators
- Learn orchestration policies from logs using imitation learning and reinforcement learning
- Improve performance-cost trade-offs in long-horizon tasks
- Extend the research toward safety and security in agent decision-making systems